

Contents

Introduction	1
Introduction	1
Materials and Methods	5
Hardware	5
Data	6
Workflow	6
Plans	7
Error simulation	8
Scoring	9
Results	10
Six open-weight implementers reach frontier accuracy on the most detailed Opus-authored plan	10
Plan detail is the binding constraint; the cliff and its single-line repair	11
qwen3.6:27b is the only fully-open Apache-2.0 implementer that does not trip on the cliff . .	11
Cross-platform: same accuracy, different wall time	12
Defensive scripting under error injection: qwen3.6:27b matches Opus on every platform . . .	12
Coda: Sonnet and Haiku confirm the implementer role for the full Anthropic family	13
References	13
Supplement	16
Plan files	16
Plan v1 (lean) — <code>plan/PLAN_v1.md</code>	16
Plan v1.25 — <code>plan/PLAN_v1p25.md</code>	17
Plan v1.5 — <code>plan/PLAN_v1p5.md</code>	17
Plan v1g — <code>plan/PLAN_v1g.md</code>	19
Plan v2 (detailed) — <code>plan/PLAN.md</code>	20
Plan v2_defensive — <code>plan/PLAN_v2_defensive.md</code>	22

Introduction

Introduction

Frontier models from Anthropic, OpenAI, and Google write quality working code for data analysis tasks, but they cost cents to dollars per call, resulting in quickly ballooning bills. Asking Opus to write a hundred-line bash script every time a postdoc aligns a few FASTQ files spends the budget that should go to harder problems. We ask a simple question: can a frontier model write the recipe once, and a free, small open-weight model running on the lab’s own hardware turn that recipe into a working script every time after?

Several groups have already asked what large language models (LLMs) can do for bioinformatics, but the questions they pose differ from ours. The earliest evaluations measured single-shot recall on genomics question-and-answer sets: GeneGPT [22] gave a closed model live access to NCBI web APIs and reported strong scores on the GeneTuring benchmark, and the 2025 GeneTuring re-run [23] put 1,600 questions to sixteen model configurations spanning closed frontiers (GPT-4o, Claude 3.5, Gemini) and small biomedical models (BioGPT, BioMedLM). A second strand swapped trivia for code on isolated puzzles. BioCoder [24] graded 2,269 short bioinformatics functions against unit tests and ran an open-versus-closed bake-off, but with code-completion models (InCoder, SantaCoder, StarCoder, CodeGen) that pre-date current open-weight chat models (publicly released model weights that anyone can download and run locally). BioLLMBench [25] scored GPT-4, Gemini and LLaMA on 24 tasks and found LLaMA often failed to emit runnable code, while an out-of-the-box study of 104 Rosalind problems [26] put GPT-3.5 ahead of GPT-4o and Llama-3-70B and compared all three to human solvers. More recent work moves from snippets to whole pipelines and to agentic systems (LLMs that plan, call tools, and revise their own output in a loop): BixBench [27] scores closed models on 50+ Dockerised computational-biology scenarios at roughly 17 % open-answer accuracy; BioMaster [28] wraps a Plan/Task/Debug/Check loop around retrieval-augmented generation (RAG, where the model is fed relevant documentation at query time) for RNA-seq, ChIP-seq, scRNA-seq and Hi-C; BioAgents [29]

fine-tunes small open-weight models for local execution and claims expert-level performance on conceptual tasks; and two workflow-language studies generate Galaxy and Nextflow pipelines with closed models plus DeepSeek-V3 [30] or rerank Galaxy workflows with Mistral-7B against GPT-4o [31].

We differ from this body of work in four ways. First, no prior study separates the **recipe** (the natural-language plan a frontier model writes once) from the **implementer** (the local model that turns that recipe into bash on every run); this is the split that lets a lab pay for cloud inference once and then work offline, and the experimental architecture we test here. Second, accuracy is almost always measured as a string match on a Q-and-A item or a unit test on an isolated function, not as whether a full pipeline runs to a VCF on real FASTQ input. Third, when open-weight models are tested at all, the hardware is left implicit; we report the same task on a Jetson AGX Orin, an RTX 5080 desktop, two Apple-silicon MacBooks and a 2× RTX A5000 workstation, so a reader can map model choice to the box on their bench. Fourth, every prior end-to-end evaluation either tests closed frontier models alone [22, 23, 27, 28, 30] or pits one or two older open-weight models against them [24, 25, 26, 31]; we instead test the latest 2026-release open-weight models against the same task. Together these gaps motivate the design of the present study: **claude-opus-4-7** writes a series of plans of increasing detail; six 2026-release open-weight models implement those plans on a desktop GPU; the strongest implementer is then cross-platform validated on three additional machines and stress-tested with a programmatic error-injection harness.

To decide which open models to use we first need to survey the landscape of available models. For 2026 this landscape is summarized in Table 1 below.

Table 1. Major open-weight model families as of May 2026.

Lab	Family / latest	Sizes (total; active for MoE)	Architecture	Coder variant	Reasoning variant	License
Meta	Llama 4 [1]	Scout 109 B (17 B-A); Maverick 400 B (17 B-A); Behemoth ~2 T (unreleased)	Sparse MoE; multimodal	—	—	Llama 4 Community License (700 M-MAU cap; excludes EU users)
Alibaba	Qwen3.6 [2]	0.6 B → 397 B (17 B-A); 27 B dense	Dense + MoE; thinking toggle	Qwen3-Coder-Next	Qwen3 thinking mode	Apache 2.0
DeepSeek	DeepSeek-V4 [3]	V4-Flash 284 B (13 B-A); V4-Pro 1.6 T (49 B-A)	Sparse MoE	DeepSeek-Coder	V4-Pro Max mode	MIT
Mistral	Mistral Large 3, Mixtral, Codestral, Ministral [4]	3 B → 675 B (41 B-A)	Granular MoE; flagship; dense small	Codestral 25.08	—	Apache 2.0 (Codestral non-production)
Google	Gemma 4 [5]	E2B, E4B, 26 B-MoE, 31 B-dense	Dense + MoE; multimodal	—	—	Apache 2.0
IBM	Granite 4.1 / 4.0 [6]	350 M, 1 B, 3 B, 8 B, 30 B; 4.0-H-Small 32 B (9 B-A)	Dense (4.1); hybrid Mamba-2 / MoE (4.0)	Granite-Code	—	Apache 2.0

Lab	Family / latest	Sizes (total; active for MoE)	Architecture	Coder variant	Reasoning variant	License
Allen AI	OLMo 3 [7]	7 B, 32 B	Dense	—	OLMo 3-Think	Apache 2.0; fully open (weights + data + recipes)
Microsoft	Phi-4 [8]	mini 3.8 B, multimodal 5.6 B, 14 B	Dense; multimodal	—	Phi-4-reasoning 14 B	MIT
NVIDIA	Nemotron-3 (<i>new in late 2025</i>) [9]	Nano-Omni 30 B (3 B-A); Super 120 B (12 B-A)	Hybrid Mamba-Transformer MoE	—	built-in agentic stack	NVIDIA Open Model License; data + recipes
Cohere	Command A, Aya [10]	Aya 3.35 B → Command A 111 B	Dense	—	—	CC-BY-NC 4.0 — research only

Legend.

- **Architecture.** Modern language models read input one piece at a time (each piece, called a *token*, is roughly a short word or word-fragment) and predict the next piece, by passing each token through a deep stack of trainable mathematical operations. Three architectural styles appear in the table:
 - *Dense.* Every internal weight is used on every token — the original transformer design and what most pre-2024 LLMs were. Simple to deploy, but compute per token scales with the model’s full size.
 - *Sparse Mixture-of-Experts (MoE).* The model is split into many parallel “expert” sub-networks plus a small “router” that, for each token, picks one or two experts to handle it; the other experts stay idle. This lets the model be very large in *total* parameters (high stored capacity) while spending the per-token compute of a much smaller model. The cost is memory: every expert still has to be loaded, even the ones that won’t fire on a given input. Every flagship released since late 2025 uses this design.
 - *Hybrid Mamba-Transformer.* A newer variant that replaces some standard layers with *state-space* (Mamba) layers. The practical consequence is that compute on a long input grows linearly with input length rather than quadratically — relevant now that flagships accept inputs of 500,000 to 1,000,000 tokens (a small book).
 - *Multimodal.* Accepts images, and sometimes audio, alongside text.
- **Parameters (B = billion, T = trillion).** Parameters are the trained numerical weights inside a model — loosely analogous to synapse strengths. More parameters mean more stored capacity but also more memory and more arithmetic per token. Two numbers matter:
 - *Total* parameters set disk and GPU-memory cost. A 70 B-parameter model occupies roughly 35–40 GB of memory at the standard 4-bit compression used for inference (a 16-bit “full precision” version would need ~140 GB).
 - *Active* parameters (notation 17 B-A = 17 billion active per token) set the per-token compute cost. For dense models the two numbers are equal. For MoE models the active count is much smaller than the total because most experts don’t fire on a given token. A 109 B (17 B-A) MoE runs about as fast per token as a 17 B dense model, but the machine still has to keep all 109 B parameters loaded.
- **Coder variant.** A general model further trained on a large corpus of source code (and tested on standard code-task benchmarks such as SWE-bench and HumanEval). Coder variants beat their general siblings on bash, Python, and refactor tasks but lose ground on chat, math, and translation.

- **Reasoning variant.** A model that, before answering, writes out a step-by-step reasoning trace (“chain of thought”) that the user typically does not see; training rewards the model when its final answer is verifiably correct. Pioneered by DeepSeek-R1 (January 2025); every major lab has since copied the recipe. Reasoning variants score higher on math, code, and multi-step problems but take roughly $5\times\text{--}50\times$ longer per call (and cost $5\times\text{--}50\times$ more).

Table 1 has notable absences: Apple’s 3-billion-parameter on-device model ships only inside iOS and macOS 26, and xAI has openly released only Grok-1 (March 2024) and Grok-2.5 (August 2025) — everything newer is closed. Every flagship since late 2025 is a sparse MoE, but dense persists below ~ 40 B because it is simpler to deploy, and Alibaba’s April 2026 27 B dense Qwen3.6 beats their own 397 B MoE on coding benchmarks — training quality, not parameter count, now sets the ceiling.

Two practical filters narrow the choice for most labs. First, license: most weights now ship under Apache 2.0 or MIT (DeepSeek, Qwen, Gemma 4, Mistral Large 3, Granite, OLMo, Phi), but Meta’s Llama 4 retains a custom Community License with a 700-million-monthly-active-user cap and an explicit exclusion of EU users, and Cohere’s Command and Aya are research-and-non-commercial only. Anyone publishing a method built on a given model has to know which license applies. Second, hardware: the four size tiers in Table 1 map to four very different machines, from a \$400–\$600 consumer card for the smallest models to a multi-GPU server costing several hundred thousand dollars for the trillion-parameter tier (Table 2). Allen AI’s OLMo and NVIDIA’s Nemotron also release the training data and recipes, not just the weights — relevant when independent replication, not just inference, matters.

Table 2. GPU options and ballpark May 2026 US street prices for running each Table 1 model class locally at 4-bit compression. Where multiple cards are needed, the listed price is the per-card cost.

Model class (Table 1 example)	GPU memory needed	GPU options (May 2026 USD)	Refs
7–8 B dense (Phi-4, Granite 8B, OLMo 3 7B)	$\sim 5\text{--}6$ GB	NVIDIA RTX 5060 Ti 16 GB ($\sim \$560$); RTX 4060 Ti 16 GB ($\sim \430); RTX 5070 12 GB ($\sim \$635$); Apple M4 Mac mini 16 GB unified ($\sim \$600$)	[11, 16]
27–32 B dense (Qwen3.6-27B, Gemma 4 31B, OLMo 3 32B)	$\sim 17\text{--}22$ GB	NVIDIA RTX 5090 32 GB ($\sim \$2,900\text{--}\$3,500$, street price 50–75 % over \$1,999 MSRP); RTX 4090 24 GB ($\sim \$1,500\text{--}\$2,200$, EOL Oct 2024); RTX A5000 24 GB ($\sim \$700\text{--}\$1,400$ used); Apple Mac Studio M3 Ultra 96 GB unified ($\sim \$4,000$)	[11, 12, 16]

Model class (Table 1 example)	GPU memory needed	GPU options (May 2026 USD)	Refs
100–400 B MoE (Llama 4 Scout 109B, Maverick 400B)	~55–250 GB	NVIDIA RTX Pro 6000 Blackwell 96 GB (~\$8,500); 2× RTX 5090 (~\$6,000–\$7,000 total); RTX 6000 Ada 48 GB (~\$6,800); H100 80 GB (~\$25K–\$33K); AMD Instinct MI300X 192 GB (~\$15K–\$20K, OEM-only); Apple Mac Studio M3 Ultra 256 GB unified (~\$9,500)	[11, 12, 13, 14, 16]
Trillion-parameter MoE (DeepSeek-V4-Pro 1.6T)	~700+ GB	NVIDIA H200 141 GB (~\$31K–\$40K each); B200 192 GB (~\$35K–\$55K each); 8-GPU DGX H200 server (~\$350K–\$500K); GB200 NVL72 rack (~\$3M+); cloud rental on B200 ~\$2.25–\$16/GPU-hour	[13, 15]

May 2026 prices are dominated by an ongoing HBM/GDDR7 shortage; consumer NVIDIA 50-series cards and high-RAM Apple Mac Studio configurations sit 30–75 % above launch MSRP. AMD MI300X and the newer MI325X are sold almost exclusively through OEM channels — public per-card numbers reflect bulk pricing or cloud rental rates [13]. The used market is liquid for retired flagships (RTX 4090, RTX A5000) and is often the cheapest path into the 27–32 B tier [11].

Given this (rapidly evolving) landscape we decided to do the following experiment: take a common sequencing data processing workflow and ask open models running on hardware accessible to an average research lab to design and execute the analysis. In doing so we experimented with a range of possibilities ranging from allowing open models to figure out everything by themselves to guiding them using a very detailed plan produced by commercial frontier models. We further complicated these tasks by simulating a variety of errors that may occur during workflow execution.

Materials and Methods

Hardware

We selected five different computers listed in Table 3. It is a combination ranging from an old workstation saved from university salvage and a desktop with a gaming GPU to the latest MacBooks and a purpose-built inexpensive NVIDIA device—NVIDIA Jetson AGX Orin. The Orin is a “RaspberryPi”-like offering from NVIDIA that costs under \$2,000 and has a very small footprint making it an idea lab-ready tiny but powerful workstation.

Table 3. Test machines used in this study.

Computer	Manufacturer	Year released	RAM	OS	GPU
NVIDIA Jetson AGX Orin Developer Kit	NVIDIA	2022	64 GB LPDDR5 (unified with GPU)	Ubuntu 22.04 LTS (NVIDIA JetPack 6)	Integrated Ampere, 2,048 CUDA cores + 64 Tensor cores
RTX 5080 desktop	Dell	2025 (GPU)	64 GB DDR5 (system)	Linux (Ubuntu)	NVIDIA RTX 5080, 10,752 CUDA cores, 16 GB GDDR7
MacBook Air M4 (24 GB)	Apple	2025	24 GB LPDDR5X (unified with GPU)	macOS 26	Apple M4 integrated GPU, 10 cores
MacBook Pro M4 Pro (48 GB)	Apple	2024	48 GB LPDDR5X (unified with GPU)	macOS Sequoia 15.6	Apple M4 Pro integrated GPU, 16 to 20 cores
2× NVIDIA RTX A5000 desktop	Dell	2021 (GPUs)	128 GB DDR4 (system)	Linux (Ubuntu 25.10)	2× NVIDIA RTX A5000, 8,192 CUDA cores each, 24 GB GDDR6 each (48 GB total)

System RAM listed for the two Linux desktops reflects the build configuration; for inference workloads the relevant memory is the GPU VRAM (last column). For the Jetson and the MacBook, RAM is unified between CPU and GPU and the model can use up to roughly the listed RAM minus the operating-system reservation.

A model’s GPU-memory footprint scales with its parameter count: at the 4-bit quantization used for all local runs here, roughly 0.5 GB per billion parameters. The plan-gradient sweep that anchors this manuscript runs on the **RTX 5080 desktop** with six 2026-release open-weight implementers — `gemma4:26b`, `gemma4:e4b`, `glm-4.7-flash`, `gpt-oss:20b`, `qwen3.6:27b`, and `qwen3.6:35b-a3b` — chosen so the family covers small dense (4 B effective), mid-size dense (26–27 B), and small-active sparse-MoE architectures within or just above the 5080’s 16 GB VRAM ceiling at 4-bit. The strongest implementer (`qwen3.6:27b`, 17 GB at 4-bit) is then cross-platform validated on the Jetson AGX Orin (64 GB unified memory), MacBook Pro M4 Pro (48 GB unified memory), and 2× RTX A5000 (48 GB total VRAM) — all three host the model comfortably. The Anthropic API models (`claude-opus-4-7` as plan author throughout, plus `claude-sonnet-4-6` and `claude-haiku-4-5` as confirmation implementers) run remotely and are independent of host GPU.

Data


We selected a small dataset [17] derived from our previous work on the analysis of mutational patterns in human mitochondria [18]. It contains four paired-end Illumina samples derived from blood and cheek tissues of a mother–child pair. It contains two fixed changes and a low-frequency variant in the child’s cheek sample — a heteroplasmy example.

Workflow

We developed a simplified version of a haploid variant-calling workflow (original: [19]) that omits pre-processing steps before variant calling and the variant-annotation phase [20, 21]. The structure of the workflow is shown in Fig. 1 below.

Figure 1. Data flow of the simplified mtDNA variant-calling workflow used in this study. Per-sample steps

(alignment through tabix) run independently for each of the four samples; the only inter-sample dependency is the final `bcftools query + awk` fan-in that builds `collapsed.tsv`.



Plans

We wanted to ask free models to perform variant calling using these data in several ways: with no plan, and guided by several plans of increasing granularity. The plans are briefly described in Table 4 below and listed in their full form in the Supplement section. Every plan was authored by `claude-opus-4-7` — the planner role in the recipe-implementer split. Opus wrote v1 (the lean reference); v2 (a hyper-detailed recipe written in response to v1’s poor implementer performance, reported in Results); v1.25 and v1.5 (controls that localize what part of v2 is doing the work); v1g (a robustness check authored mechanically from a public Galaxy tool registry rather than by Opus); v0.5 (a no-recipe condition with only the tool order); and v2_defensive (v2 plus an error-handling skeleton).

Each model is run under one of two conditions, which we call **tracks**. In **Track A**, the model receives a written plan (one of the rows of Table 4 with a non-empty *File* column) together with the tool inventory and writes `run.sh` from the plan. In **Track B**, the model receives the problem statement and the tool inventory only — no plan — and writes `run.sh` from scratch. Track B exists to bound how much of the implementer’s behaviour comes from the model’s prior training versus from the supplied recipe; every other plan in Table 4 is a Track A condition.

Table 4. Plan variants used in this study. Plans of increasing granularity, plus two “no-plan” controls. File sizes are bytes; full text for each plan file is in the Supplement.

Plan	File	Size	Short summary	Hypothesis tested
Track B	<i>(no plan)</i>	0	Problem statement + tool inventory only; no implementation hints	How much can a model do from scratch?
v0.5	<code>prompts/track_b_with_order_user.txt</code>	1.4 KB	Track B + one line giving the tool order: <code>bwa → samtools → lofreq → bcftools → awk</code>	Does sequencing alone help local models?
v1 (lean)	<code>plan/PLAN_v1.md</code>	3.1 KB	Numbered bullets naming tools and key flags; no exact command lines	Reference lean plan
v1.25	<code>plan/PLAN_v1p25.md</code>	3.1 KB	v1 + the exact <code>lofreq call-parallel</code> command line	Does one full command bridge the v1 → v2 gap?

Plan	File	Size	Short summary	Hypothesis tested
v1.5	plan/PLAN_v1p5.md	1.3 KB	v2 with every prose paragraph and Gotchas block stripped — code fences only	Are the prose explanations load-bearing or decorative?
v1g	plan/PLAN_v1g.md	4.2 KB	v1 + a Galaxy-IUC-mechanical lofreq snippet (extracted from <code>tools-iuc@39e7456</code>)	Can a tool registry replace a human plan author?
v2 (detailed)	plan/PLAN.md	4.6 KB	Every step gives the exact command line plus Gotchas block	Reference detailed plan
v2_defensive	plan/PLAN_v2_defensive.md	6.5 KB	v2 + <code>try()</code> helper, output validation after every step, retry-once, per-sample isolation, structured failure log	Does explicit error-handling prose make implementer scripts defensive against runtime tool failures?

Error simulation

To test how each model handles tool failures, we wrapped `bwa` and `lofreq` in short shell scripts (“shims”) that imitate the real binaries on disk. Before each run, the harness writes the two shims into a per-run directory and adds that directory to the front of `PATH` — the colon-separated list of directories the shell consults when it needs to find an executable. When the script types `bwa mem`, the shell finds the shim first and runs it instead of the real `bwa`. Each shim reads two environment variables: `EVAL_INJECT_PATTERN` names the failure type and `EVAL_INJECT_TARGET` names which of `bwa` or `lofreq` the failure applies to. For every untargeted invocation the shim passes the call straight through to the real binary; for the targeted invocation it injects the failure and then either passes through or stops. The model’s `run.sh` does not know it is being intercepted.

Seven failure patterns cover the kinds of breakage a real lab pipeline encounters (Table 5). Each pattern is paired with one tool (or both), so a full error-matrix cell is one (model $\times$ plan $\times$ pattern $\times$ target tool $\times$ seed) combination. We use seeds 42, 43 and 44, and run every cell on both the **v2** and the **v2_defensive** plan to measure how the defensive prose changes outcomes.

Table 5. Failure patterns injected by the PATH shims into `bwa` and / or `lofreq`.

Pattern	Affected tool	What the shim does
<code>flake_first_call</code>	both	First call exits 1; later calls pass through.
<code>one_sample_fails</code>	both	Shim exits 1 only when <code>M117C1-ch</code> appears in the command line; the other three samples run clean.

Pattern	Affected tool	What the shim does
<code>slow_tool</code>	both	Sleep 30 s, then run the real tool — output identical, just delayed.
<code>stderr_warning_storm</code>	both	Dump 200 lines of <code>WARNING:</code> text to <code>stderr</code> , then run the real tool — output is correct but the <code>stderr</code> looks alarming.
<code>missing_lib_error</code>	both	Exit 127 without invoking the real tool, mimicking a missing shared library (<code>error while loading shared libraries: libhts.so.3: cannot open shared object file</code>).
<code>silent_truncation</code>	lofreq only	Run the real <code>lofreq</code> , then truncate the <code>-o</code> output file to zero bytes — <code>lofreq</code> returns 0 but the VCF is empty.
<code>wrong_format_output</code>	lofreq only	Run the real <code>lofreq</code> , then strip every variant line from the VCF; only the header survives — the file still parses, but contains no calls.

Beyond the variant-overlap score M3, we score each cell on three error-handling metrics. `m_handle` classifies the script’s overall response as `crash` (exited non-zero with no failure log), `propagate` (the failure passed through to a downstream step), `partial` (the script aborted but recorded the failure structurally), or `recover` (the script completed successfully despite the injection). `m_recover` is binary: did the script produce the count of valid VCFs that the failure pattern allows? For `one_sample_fails` the best achievable is 3 (the bad sample’s VCF is correctly absent); for `silent_truncation`, `wrong_format_output` and `missing_lib_error` the best is 0 (the script should detect the broken or missing output and skip it); for the remaining three patterns the best is 4. `m_diagnose` is binary: did the script announce the failure — through a populated `failures.log`, a summary line on `stderr`, or a sample-name-and-failure-word pair in its output — or did it stay silent? All three metrics are defined in `score/score_run.py:error_handling`.

Scoring

We score every run on five metrics computed automatically by `score/score_run.py` from the script’s outputs and the canonical answer key. The metrics range from coarse (“did anything execute?”) to fine (“which exact variants did we call, and at what allele frequencies?”) so that we can tell apart a run that produced no output, a run that produced output of the wrong content, and a run that got every call right.

Table 6. Per-run scoring metrics.

Metric	Type	What it measures	Pass condition
M1 (executes)	binary	Did the script run to completion?	<code>bash run.sh</code> returns exit 0 within the 600-second wall-clock budget.
M2 (schema)	binary	Did the script produce the right <i>files</i> ?	All nine expected outputs exist (4 BAM + 4 BAI + 4 VCF.gz + 4 TBI + <code>collapsed.tsv</code>) and <code>bcftools view -h</code> parses each VCF header.

Metric	Type	What it measures	Pass condition
M3 (variant overlap)	float in $[0, 1]$	Did the script call the right <i>variants</i> ?	Per sample, the Jaccard overlap of (<code>chrom</code> , <code>pos</code> , <code>ref</code> , <code>alt</code>) between the script’s PASS-or-unfiltered records and the answer key, requiring matched allele frequencies within ± 0.02 . Averaged over the four samples.
M4 (cost and time)	quantitative	What did the run cost?	Tokens consumed; USD spent (Anthropic only — Ollama is free); generation time and execution time, both in seconds.
M5 (script quality)	binary	Is the script defensible bash?	All four of: <code>set -euo pipefail</code> present; no absolute user paths (<code>/home/anton/...</code>); <code>shellcheck</code> clean; idempotent (a second run on a fully populated output directory exits 0 without redoing work).

M3 is the headline score in every results section. It is the metric that captures whether the model has actually done the job, and it is the only one that distinguishes a script that produced four well-formed VCFs of the wrong content from one that produced four correct VCFs. A run that fails M1 (the script did not execute to completion) gets $M2 = 0$ by construction; we still compute M3 over whatever VCFs the script did produce and treat any missing file as zero overlap on its sample. A model that uses a different but valid tool from the one named in the plan — `bcftools mpileup` in place of `lofreq`, for example — gets credit on M3 for the variants it called, not for the tool it chose.

For the error-injection cells (described under *Error simulation*) we add the three further metrics — `m_handle`, `m_recover`, `m_diagnose` — that score how the script behaved when one of its tools deliberately misbehaved. All eight metrics are written to a per-run `score.json` and aggregated by `score/aggregate.py` into the table used to produce every figure and table in the Results section.

Results

Six open-weight implementers reach frontier accuracy on the most detailed Opus-authored plan

We ran six 2026-release open-weight implementers on the RTX 5080 against every plan in Table 4 (Track A and Track B), three seeds per cell, scoring M3 — the share of called variants matching the truth set within a 0.02 allele-frequency window (Methods, *Scoring*). On the v2 plan, the most detailed Opus-authored recipe, five of the six implementers reach $M3 = 1.000$ across all three seeds (`qwen3.6:27b`, `qwen3.6:35b-a3b`, `gemma4:26b`, `gemma4:e4b`, `glm-4.7-flash`); the sixth, `gpt-oss:20b`, reaches mean $M3 = 0.67$, a budget effect — its built-in reasoning mode consumes most of the output token allowance before the script is finished. The recipe-implementer split is operational on commodity hardware: a frontier-tier author (Opus) writes the recipe once; a 17 GB open-weight implementer reproduces frontier accuracy at zero per-call cost on a desktop GPU.

Plan detail is the binding constraint; the cliff and its single-line repair

We constructed plans of increasing detail to localize what an implementer needs from the recipe (Figure 2). At the lean end, `claude-opus-4-7` first wrote `v1`: a numbered bullet list naming each tool and its key flags but no exact command lines. `v1` produces a clean cliff: `qwen3.6:27b` and `glm-4.7-flash` reach $M3 = 1.000$, but the four other implementers crash at $M3 \approx 0$. The detailed `v2` recipe, written in response, recovers all six implementers (subject to the `gpt-oss:20b` budget effect noted above). To isolate which part of `v2` is doing the work, we authored two intermediate plans: `v1.25` = `v1` plus the single literal `lofreq call-parallel --pp-threads 4 -f data/ref/chrM.fa -o results/{s}.vcf results/{s}.bam` invocation, and `v1.5` = `v2` with every prose paragraph and *Gotchas* block stripped, code fences only. `v1.25` reproduces the `v2` jump for every dense ≥ 26 B implementer (`gemma4:26b`, `qwen3.6:27b`, `qwen3.6:35b-a3b`); `v1.5` does not. The active ingredient is the literal command syntax, not the surrounding explanation. Track B (no plan) and `v0.5` (no plan plus the tool order `bwa` $\rightarrow$ `samtools` $\rightarrow$ `lofreq` $\rightarrow$ `bcftools` $\rightarrow$ `awk`) confirm the floor: every open-weight implementer scores $M3 \approx 0$ in both no-recipe conditions; supplying tool order without flags or arguments does not move the cliff.




Figure 1: Figure 2. Mean M3 across three seeds for six open-weight implementers $\times$ seven Opus-authored plan variants, on the RTX 5080. Columns are ordered by plan detail (left = no plan, right = most detailed). Anthropic API plan-author rows are not shown here because Anthropic models execute on Anthropic’s hardware; they appear in the cross-platform error-injection panel below.

A robustness control rules out the alternative explanation that the `v2` effect is specific to Opus’s writing style: `v1g` is `v1` plus the `lofreq call-parallel` snippet extracted *mechanically* from the Galaxy IUC tool registry (`tools-iuc039e7456`), a hand-curated public corpus of bioinformatics tool wrappers. `v1g` produces a similar implementer lift to `v1.25` — the load-bearing element is the command syntax, regardless of whether a frontier model or a tool registry is the source.

qwen3.6:27b is the only fully-open Apache-2.0 implementer that does not trip on the cliff

Within the six-implementer lineup, two pass `v1` with $M3 = 1.000$: `qwen3.6:27b` (Apache 2.0, Alibaba; 17 GB at 4-bit) and `glm-4.7-flash` (Z.ai; comparable size). `qwen3.6:27b` is the cleaner choice for a working lab: an Apache-2.0 license with no field-of-use restrictions, a single-file `ollama pull qwen3.6:27b` deployment, and a memory footprint that fits the 24 GB VRAM tier on Table 2 with room. We therefore selected `qwen3.6:27b` as the protagonist for the cross-platform validation that follows.

Cross-platform: same accuracy, different wall time

qwen3.6:27b was tested on the three remaining platforms — Jetson AGX Orin, MacBook Pro M4 Pro, and 2× RTX A5000 — under the v2 plan. Mean M3 = 1.000 on every platform on every seed. Wall-clock generation time differs sharply (Table 7): the 2× A5000 returns in under half a minute, the M4 Pro and Jetson take 1.5 to 2 minutes, and the RTX 5080 takes about 5 minutes — the 5080 row reflects the model partially spilling out of its 16 GB VRAM into system RAM, which Ollama handles automatically but slowly. None of the wall times is too slow for a working lab; pick the cheapest box that fits the model in dedicated or unified memory and the M3 score is unchanged.

Table 7. **qwen3.6:27b** on the v2 plan: median wall-clock time per generation (seconds), inter-quartile range in brackets. M3 is 1.000 on every cell across every platform.

Platform	VRAM/UMA	Median wall time (s)	IQR (Q1–Q3)	n
2× RTX A5000	48 GB total VRAM (fits)	29	[24–31]	36
MacBook Pro M4 Pro	48 GB unified (fits)	92	[91–136]	36
NVIDIA Jetson AGX Orin	64 GB unified (fits)	105	[98–107]	36
RTX 5080 desktop	16 GB VRAM (spills to RAM)	302	[288–322]	6

Defensive scripting under error injection: **qwen3.6:27b** matches **Opus** on every platform

To probe whether **qwen3.6:27b** handles the *robustness* case as well as the happy path, we ran the 36-cell error-injection matrix (12 PATH-shim cells × 3 seeds; Methods, *Error simulation*) on **qwen3.6:27b** and **claude-opus-4-7** on three platforms (Jetson, M4 Pro, 2× A5000), under both the v2 baseline plan and the v2_defensive plan that adds a `try() { ... }` helper, per-step output validation, and a structured `failures.log` (Methods, *Plans*). Both models behave **identically** at the cell level on every platform tested (Figure 3, Table 8). On the v2 baseline, both produce 15 recover / 0 partial / 6 propagate / 15 crash; on v2_defensive, both produce 21 recover / 15 partial / 0 propagate / 0 crash — the defensive prose drains every crash into a structurally-recovered partial — and both score mean M3 = 0.625 on the matrix. A 17 GB open-weight implementer matches a frontier-API plan author cell-for-cell and score-for-score on the error matrix.

Table 8. Per-cell behaviour signature and mean M3 for **qwen3.6:27b** and **claude-opus-4-7** on the 36-cell error matrix (12 patterns × 3 seeds), across Jetson, M4 Pro, and 2× A5000. Both plans, both models, all three platforms summarised.

Model	Plan	Platform	recover	partial	propagate	crash	mean M3
claude-opus-4-7		Jetson, M4, A5000 (identical)	15	0	6	15	0.333
claude-opus-4-7	v2_defensive	Jetson, A5000, M4 r2 (identical)	21	15	0	0	0.625
qwen3.6:27b		Jetson, M4 r2, A5000 (identical)	15	0	6	15	0.333

Model	Plan	Platform	recover	partial	propagate	crash	mean M3
qwen3.6:27b2	v2_defensive	Jetson, M4 r2, A5000 (identical)	21	15	0	0	0.625

The error-injection result closes the loop on the recipe-implementer architecture: under both the happy-path plan and a defensive plan, a free, locally-runnable Apache-2.0 implementer handles the same workload as the frontier API author at the same accuracy.

Coda: Sonnet and Haiku confirm the implementer role for the full Anthropic family

We also ran `claude-sonnet-4-6` and `claude-haiku-4-5` against the same error matrix on the Jetson and 2× A5000 platforms. Both produce the same v2_defensive frontier signature as Opus and `qwen3.6:27b` (Sonnet: 21 recover / 14 partial / 0 propagate / 1 crash on Jetson, 21/15/0/0 on A5000, mean M3 = 0.583–0.625; Haiku: 21/15/0/0 on both, mean M3 = 0.625). The implementer role generalizes across the Anthropic family. Cross-platform timing for Anthropic models is not informative for hardware-throughput claims — the LLM call is remote — so we report Sonnet and Haiku as confirmation of the implementer role, not as a hardware comparison.

References

- [1] Meta. Llama 4: a new crop of flagship AI models. *TechCrunch*, April 5, 2025. <https://techcrunch.com/2025/04/05/meta-releases-llama-4-a-new-crop-of-flagship-ai-models/>
- [2] Alibaba (Qwen team). Qwen3.6 family. GitHub. <https://github.com/QwenLM/Qwen3.6>
- [3] DeepSeek-AI. DeepSeek-V4 preview release notes. DeepSeek API documentation, April 24, 2026. <https://api-docs.deepseek.com/news/news260424>
- [4] Mistral AI. Introducing Mistral 3. Mistral AI news, December 2, 2025. <https://mistral.ai/news/mistral-3>
- [5] Google. Introducing Gemma 4. Google blog, April 2, 2026. <https://blog.google/innovation-and-ai/technology/developers-tools/gemma-4/>
- [6] IBM Research. Granite 4.1 AI foundation models. April 30, 2026. <https://research.ibm.com/blog/granite-4-1-ai-foundation-models>
- [7] Allen Institute for AI (Ai2). OLMo 3. November 20, 2025. <https://allenai.org/blog/olmo3>
- [8] Microsoft. Welcome to the new Phi-4 models — Phi-4-mini and Phi-4-multimodal. Microsoft TechCommunity, Educator Developer Blog. <https://techcommunity.microsoft.com/blog/educatordeveloperblog/welcome-to-the-new-phi-4-models—microsoft-phi-4-mini-phi-4-multimodal/4386037>
- [9] NVIDIA. NVIDIA debuts Nemotron 3 family of open models. NVIDIA Newsroom. <https://nvidianews.nvidia.com/news/nvidia-debuts-nemotron-3-family-of-open-models>
- [10] Cohere. Models — Command A, Aya. Cohere docs. <https://docs.cohere.com/docs/models>
- [11] BestValueGPU. Consumer NVIDIA RTX GPU price history and specifications (RTX 4060 Ti, 4090, 5060 Ti, 5070, 5090). <https://bestvaluegpu.com/>
- [12] Thunder Compute. NVIDIA RTX Pro 6000 Blackwell pricing analysis. <https://www.thundercompute.com/blog/nvidia-rtx-pro-6000-pricing>
- [13] Thunder Compute. AMD Instinct MI300X pricing. <https://www.thundercompute.com/blog/amd-mi300x-pricing>

qwen3.6:27b matches claude-opus-4-7 across platforms on the error matrix




Figure 2: Figure 3. Per-cell error-handling outcome (recover / partial / propagate / crash) for claude-opus-4-7 and qwen3.6:27b across the 12-cell PATH-shim error matrix, on three hardware platforms. Left column = v2 baseline plan; right column = v2_defensive plan. Rows within each panel are the two models. Cells are color-coded by the modal `m_handle` outcome across three seeds.

- [14] Jarvis Labs. NVIDIA H100 80 GB pricing guide. <https://jarvislabs.ai/blog/h100-price>
- [15] Northflank. NVIDIA B200 cost analysis and cloud rental rates. <https://northflank.com/blog/how-much-does-an-nvidia-b200-gpu-cost>
- [16] Apple. Mac Studio configurations and pricing. <https://www.apple.com/mac-studio/specs/>
- [17] Nekrutenko A. Datasets for Galaxy Collection Operations Tutorial. *Zenodo* dataset, 2021. doi:10.5281/zenodo.5119008. <https://zenodo.org/records/5119008>
- [18] Rebolledo-Jaramillo B, Su MS, Stoler N, McElhoe JA, Dickins B, Blankenberg D, Korneliussen TS, Chiaramonte F, Nielsen R, Holland MM, Paul IM, Nekrutenko A, Makova KD. Maternal age effect and severe germ-line bottleneck in the inheritance of human mitochondrial DNA. *Proc Natl Acad Sci U S A*. 2014;111(43):15474–15479. <https://pubmed.ncbi.nlm.nih.gov/25313049/>
- [19] Nekrutenko A. iwc-workflows/haploid-variant-calling-wgs-pe (v0.1). *Zenodo*, March 24, 2025. doi:10.5281/zenodo.15078463. <https://zenodo.org/records/15078463>
- [20] Maier W, Bray S, van den Beek M, Bouvier D, Coraor N, Miladi M, Singh B, De Argila JR, Baker D, Roach N, Gladman S, Coppens F, Martin DP, Lonie A, Grüning B, Kosakovsky Pond SL, Nekrutenko A. Ready-to-use public infrastructure for global SARS-CoV-2 monitoring. *Nat Biotechnol*. 2021;39(10):1178–1179. <https://pubmed.ncbi.nlm.nih.gov/34588690/>
- [21] Mei H, Arbehthuber B, Cremona MA, DeGiorgio M, Nekrutenko A. A high-resolution view of adaptive event dynamics in a plasmid. *Genome Biol Evol*. 2019;11(10):3022–3034. <https://pubmed.ncbi.nlm.nih.gov/31539047/>
- [22] Jin Q, Yang Y, Chen Q, Lu Z. GeneGPT: augmenting large language models with domain tools for improved access to biomedical information. *Bioinformatics*. 2024;40(2):btac075. doi:10.1093/bioinformatics/btac075. PMID:38341654.
- [23] Shang X, Liao X, Ji Z, Hou W. Benchmarking large language models for genomic knowledge with GeneTuring. *Brief Bioinform*. 2025;26(5):bbaf492. doi:10.1093/bib/bbaf492.
- [24] Tang X, Qian B, Gao R, Chen J, Chen X, Gerstein MB. BioCoder: a benchmark for bioinformatics code generation with large language models. *Bioinformatics*. 2024;40(Suppl_1):i266–i276. doi:10.1093/bioinformatics/btac230. PMID:38940140.
- [25] Sarwal V, Andreoletti G, Munteanu V, Suhodolschi A, Ciorba D, Bostan V, Dimian M, Eskin E, Wang W, Mangul S. BioLLMBench: a benchmark for large language models in bioinformatics. *bioRxiv*; 2023. doi:10.1101/2023.12.19.572483.
- [26] Rajesh V, Siwo GH. Out-of-the-box bioinformatics capabilities of large language models (LLMs). *bioRxiv*; 2025. doi:10.1101/2025.08.22.671610. PMID:40909484.
- [27] Mitchener L, Laurent JM, Andonian A, Tenmann B, Narayanan S, Wellawatte GP, White A, Sani L, Rodrigues SG. BixBench: a comprehensive benchmark for LLM-based agents in computational biology. *arXiv:2503.00096*; 2025. <https://arxiv.org/abs/2503.00096>
- [28] Su H, Long W, Zhang Y. BioMaster: multi-agent system for automated bioinformatics analysis workflow. *bioRxiv*; 2025. doi:10.1101/2025.01.23.634608.
- [29] Mehandru N, Hall AK, Melnichenko O, Dubinina Y, Tsiurlikov D, Bamman D, Alaa A, Saponas S, Malladi VS. BioAgents: bridging the gap in bioinformatics analysis with multi-agent systems. *Sci Rep*. 2025;15:39036. doi:10.1038/s41598-025-25919-z.
- [30] Alam K, Roy B. From Prompt to Pipeline: large language models for scientific workflow development in bioinformatics. *arXiv:2507.20122*; 2025. <https://arxiv.org/abs/2507.20122>
- [31] Cynthia ST, Roy B. Towards LLM-powered task-aware retrieval of scientific workflows for Galaxy. *arXiv:2511.01757*; 2025. <https://arxiv.org/abs/2511.01757>

Supplement

Plan files

The full text of each plan file referenced in Table 4. Each plan is reproduced verbatim from `plan/` in the repository and shown inside a fenced Markdown block to preserve its original heading hierarchy and code formatting.

Plan v1 (lean) — `plan/PLAN_v1.md`

Per-sample mtDNA amplicon variant-calling plan

1. ****Set globals and prepare results directory****
 - Define `THREADS=4` and the sample list: `M117-bl M117-ch M117C1-bl M117C1-ch`.
 - Create `results/` if missing. Use `set -euo pipefail`.
 - Treat every output step as idempotent: guard each artifact with an existence check (e.g. skip if `
2. ****Reference indexing (once, in `data/ref/`)****
 - `samtools faidx data/ref/chrM.fa` → produces `chrM.fa.fai`.
 - `bwa index data/ref/chrM.fa` → produces the `.amb .ann .bwt .pac .sa` set.
 - Skip both if the index files already exist.
3. ****Per-sample alignment with `bwa mem`****
 - Use `bwa mem -t 4` with the paired FASTQs `data/raw/{sample}_1.fq.gz` and `data/raw/{sample}_2.fq.gz`.
 - Pass the read group via `-R` as a single double-quoted argument containing literal backslash-t between the two sample names:
 - exact form: `-R "@RG\tID:{sample}\tSM:{sample}\tLB:{sample}\tPL:ILLUMINA"`
 - The `\t` must remain the two characters backslash and `t` - bwa parses them itself. Do NOT use `\\t`.
 - Separators between key and value are colons `:`, not `=`.
4. ****SAM → sorted BAM****
 - Pipe `bwa mem` stdout into `samtools sort -@ 4 -o results/{sample}.bam`.
 - Do NOT run `markdup` or `rmdup`: this is amplicon data where PCR duplicates are expected and biological duplicates are not.
5. ****BAM indexing****
 - `samtools index -@ 4 results/{sample}.bam` → `results/{sample}.bam.bai`.
6. ****Variant calling with `lofreq call-parallel`****
 - Use the `call-parallel` subcommand (not plain `lofreq call`) with `--pp-threads 4`.
 - Reference: `data/ref/chrM.fa`. Input: `results/{sample}.bam`.
 - Write uncompressed VCF to a temporary path (e.g. `results/{sample}.vcf`); lofreq emits plain VCF.
7. ****VCF compression and indexing****
 - Compress with `bgzip` (not `bcftools view -O z`) producing `results/{sample}.vcf.gz`.
 - Index with `tabix -p vcf results/{sample}.vcf.gz` → `results/{sample}.vcf.gz.tbi`.
 - Remove the intermediate uncompressed `.vcf`.
8. ****Collapse step → `results/collapsed.tsv`****
 - For each sample, run `bcftools query -f '{sample}\t%CHROM\t%POS\t%REF\t%ALT\t%INFO/AF\n' results/{sample}.vcf.gz`.
 - Concatenate all four samples' output.
 - Prepend a single header line `sample\tchrom\tpos\tref\talt\taf` (tab-separated).
 - Output is tab-separated, one variant per line, header on, written to `results/collapsed.tsv`. Rebuild index.
9. ****Idempotency check****
 - Final pass: re-running the script on a fully populated `results/` exits 0, performs no work, and 1

Plan v1.25 — plan/PLAN_v1p25.md

Per-sample mtDNA amplicon variant-calling plan

1. ****Set globals and prepare results directory****
 - Define `THREADS=4` and the sample list: `M117-bl M117-ch M117C1-bl M117C1-ch`.
 - Create `results/` if missing. Use `set -euo pipefail`.
 - Treat every output step as idempotent: guard each artifact with an existence check (e.g. skip if `
2. ****Reference indexing (once, in `data/ref/`)****
 - `samtools faidx data/ref/chrM.fa` → produces `chrM.fa.fai`.
 - `bwa index data/ref/chrM.fa` → produces the `.amb .ann .bwt .pac .sa` set.
 - Skip both if the index files already exist.
3. ****Per-sample alignment with `bwa mem`****
 - Use `bwa mem -t 4` with the paired FASTQs `data/raw/{sample}_1.fq.gz` and `data/raw/{sample}_2.fq.gz`.
 - Pass the read group via `-R` as a single double-quoted argument containing literal backslash-t between the two characters backslash and `t` - bwa parses them itself. Do NOT use `\\t` - exact form: `-R "@RG\\tID:{sample}\\tSM:{sample}\\tLB:{sample}\\tPL:ILLUMINA"`
 - The `\\t` must remain the two characters backslash and `t` - bwa parses them itself. Do NOT use `\\t`
 - Separators between key and value are colons `:`, not `=`.
4. ****SAM → sorted BAM****
 - Pipe `bwa mem` stdout into `samtools sort -@ 4 -o results/{sample}.bam`.
 - Do NOT run `markdup` or `rmdup`: this is amplicon data where PCR duplicates are expected and biologically correct.
5. ****BAM indexing****
 - `samtools index -@ 4 results/{sample}.bam` → `results/{sample}.bam.bai`.
6. ****Variant calling with `lofreq call-parallel`****
 - Exact command:
...
`lofreq call-parallel --pp-threads 4 -f data/ref/chrM.fa -o results/{sample}.vcf results/{sample}.bam`
...
 - Output: `results/{sample}.vcf` (uncompressed). lofreq emits plain VCF.
7. ****VCF compression and indexing****
 - Compress with `bgzip` (not `bcftools view -O z`) producing `results/{sample}.vcf.gz`.
 - Index with `tabix -p vcf results/{sample}.vcf.gz` → `results/{sample}.vcf.gz.tbi`.
 - Remove the intermediate uncompressed `.vcf`.
8. ****Collapse step → `results/collapsed.tsv`****
 - For each sample, run `bcftools query -f '{sample}\\t%CHROM\\t%POS\\t%REF\\t%ALT\\t%INFO/AF\\n' results/{sample}.vcf.gz`
 - Concatenate all four samples' output.
 - Prepend a single header line `sample\\tchrom\\tpos\\tref\\talt\\taf` (tab-separated).
 - Output is tab-separated, one variant per line, header on, written to `results/collapsed.tsv`. Rebuild index.
9. ****Idempotency check****
 - Final pass: re-running the script on a fully populated `results/` exits 0, performs no work, and 1

Plan v1.5 — plan/PLAN_v1p5.md

Implementation Plan: Per-sample mtDNA Variant Calling

Boilerplate (top of `run.sh`)

```

...
set -euo pipefail
THREADS=4
SAMPLES=("M117-bl" "M117-ch" "M117C1-bl" "M117C1-ch")
mkdir -p results
...

All per-sample steps run in `for sample in "${SAMPLES[@]}"; do ... done`.

---

## 1. Reference indexing - BWA

...
bwa index data/ref/chrM.fa
...

## 2. Reference indexing - samtools faidx

...
samtools faidx data/ref/chrM.fa
...

## 3. Per-sample alignment + sort (one pipeline)

...
bwa mem -t 4 -R "@RG\tID:{sample}\tSM:{sample}\tLB:{sample}\tPL:ILLUMINA" data/ref/chrM.fa data/raw/{sample}
...

## 4. BAM index

...
samtools index -@ 4 results/{sample}.bam
...

## 5. Variant calling - LoFreq

...
lofreq call-parallel --pp-threads 4 -f data/ref/chrM.fa -o results/{sample}.vcf results/{sample}.bam
...

## 6. VCF compression + tabix index

...
bgzip -f results/{sample}.vcf
...
...
tabix -p vcf results/{sample}.vcf.gz
...

## 7. Collapsed TSV

...
printf 'sample\tchrom\tpos\tref\talt\tmaf\n' > results/collapsed.tsv

```

```

...
...
bcftools query -f '%CHROM\t%POS\t%REF\t%ALT\t%INFO/AF\n' results/{sample}.vcf.gz | awk -v s={sample} 'B
...

```

Plan v1g — plan/PLAN_v1g.md

Per-sample mtDNA amplicon variant-calling plan

```
<!--
```

v1g = v1 with Galaxy-IUC-derived CLI snippets injected per step where IUC has clean coverage. Extracted mechanically by scripts/galaxy_to_snippet.py from tools-iuc commit 39e745658a6ff7f013788871916574117a0f47f1 (2026-04-27).

IUC coverage map:

```

bwa, bwa-mem      : extraction yields mostly noise (heavy macro use) - fallback to v1 prose
samtools_faidx    : all-conditional command block - fallback to v1 prose
samtools_sort     : partial extraction with placeholders - fallback to v1 prose
samtools_index    : not in IUC - fallback to v1 prose
lofreq_call_parallel : clean extraction - INJECTED (step 6)
bgzip / tabix     : not in IUC - fallback to v1 prose
bcftools_query    : format string in stripped Cheetah var - fallback to v1 prose
-->

```

1. **Set globals and prepare results directory**

- Define `THREADS=4` and the sample list: `M117-bl M117-ch M117C1-bl M117C1-ch`.
- Create `results/` if missing. Use `set -euo pipefail`.
- Treat every output step as idempotent: guard each artifact with an existence check (e.g. skip if `

2. **Reference indexing (once, in `data/ref/`)**

- `samtools faidx data/ref/chrM.fa` → produces `chrM.fa.fai`.
- `bwa index data/ref/chrM.fa` → produces the `.amb .ann .bwt .pac .sa` set.
- Skip both if the index files already exist.

3. **Per-sample alignment with `bwa mem`**

- Use `bwa mem -t 4` with the paired FASTQs `data/raw/{sample}_1.fq.gz` and `data/raw/{sample}_2.fq.gz`.
- Pass the read group via `-R` as a single double-quoted argument containing literal backslash-t between exact form: `-R "@RG\tID:{sample}\tSM:{sample}\tLB:{sample}\tPL:ILLUMINA"`
- The `\t` must remain the two characters backslash and `t` - bwa parses them itself. Do NOT use `\\t`.
- Separators between key and value are colons `:`, not `=`.

4. **SAM → sorted BAM**

- Pipe `bwa mem` stdout into `samtools sort -@ 4 -o results/{sample}.bam`.
- Do NOT run `markdup` or `rmdup`: this is amplicon data where PCR duplicates are expected and biolo

5. **BAM indexing**

- `samtools index -@ 4 results/{sample}.bam` → `results/{sample}.bam.bai`.

6. **Variant calling with `lofreq call-parallel`**

- Galaxy IUC canonical invocation (extracted from `tools/lofreq/lofreq_call.xml` @ tools-iuc 39e745658a6ff7f013788871916574117a0f47f1)
- ...
- `lofreq call-parallel --pp-threads 4 --verbose`
- `--ref data/ref/chrM.fa --out results/{sample}.vcf`
- `--sig`
- `--bonf`

```
results/{sample}.bam
...
```

(The bare `--sig` and `--bonf` lines come from Galaxy-runtime-supplied values; you can omit them a

7. **VCF compression and indexing**

- Compress with `bgzip` (not `bcftools view -O z`) producing `results/{sample}.vcf.gz`.
- Index with `tabix -p vcf results/{sample}.vcf.gz` → `results/{sample}.vcf.gz.tbi`.
- Remove the intermediate uncompressed `.vcf`.

8. **Collapse step → `results/collapsed.tsv`**

- For each sample, run `bcftools query -f '{sample}\t%CHROM\t%POS\t%REF\t%ALT\t%INFO/AF\n' results/{`
- Concatenate all four samples' output.
- Prepend a single header line `sample\tchrom\tpos\tref\talt\taf` (tab-separated).
- Output is tab-separated, one variant per line, header on, written to `results/collapsed.tsv`. Rebu

9. **Idempotency check**

- Final pass: re-running the script on a fully populated `results/` exits 0, performs no work, and l

Plan v2 (detailed) — plan/PLAN.md

Implementation Plan: Per-sample mtDNA Variant Calling

Boilerplate (top of `run.sh`)

- First line after shebang: `set -euo pipefail`.
- Constants: `THREADS=4` and `SAMPLES=("M117-bl" "M117-ch" "M117C1-bl" "M117C1-ch")`.
- Create output dir: `mkdir -p results`.
- All per-sample steps must be wrapped in `for sample in "${SAMPLES[@]}; do ... done`.

```
---
```

1. Reference indexing - BWA

```
...
```

```
bwa index data/ref/chrM.fa
...
```

- Outputs (5 sibling files): `data/ref/chrM.fa.amb`, `.ann`, `.bwt`, `.pac`, `.sa`.
- Idempotency guard: `[[ -f data/ref/chrM.fa.bwt ]] || bwa index data/ref/chrM.fa`
- Gotcha: `bwa index` writes outputs next to the input; the dir must be writable. No flags needed for a

2. Reference indexing - samtools faidx

```
...
```

```
samtools faidx data/ref/chrM.fa
...
```

- Output: `data/ref/chrM.fa.fai`.
- Guard: `[[ -f data/ref/chrM.fa.fai ]] || samtools faidx data/ref/chrM.fa`

3. Per-sample alignment + sort (one pipeline)

```
...
```

```
bwa mem -t 4 -R "@RG\tID:{sample}\tSM:{sample}\tLB:{sample}\tPL:ILLUMINA" data/ref/chrM.fa data/raw/{sa}
...
```

- Output: ``results/{sample}.bam``.
- Guard: ``[[ -f results/{sample}.bam ]] || { bwa mem ... | samtools sort ... ; }`` - wrap the whole pipe
- RG string gotchas (CRITICAL):
 - Use colons (``ID:``, ``SM:``, ``LB:``, ``PL:``) - never ``=``.
 - Use the ****literal two characters**** ``\t`` (backslash + t) inside the double-quoted string. Do NOT use
 - The whole ``-R`` value must be a single double-quoted argument.
- ``samtools sort`` trailing ``-`` reads from stdin.

4. BAM index

```
...
samtools index -@ 4 results/{sample}.bam
...
```

- Output: ``results/{sample}.bam.bai``.
- Guard: ``[[ -f results/{sample}.bam.bai ]] || samtools index -@ 4 results/{sample}.bam``
- Do NOT run ``markdup`` - this is amplicon data; PCR duplicates are expected and informative.

5. Variant calling - LoFreq

```
...
lofreq call-parallel --pp-threads 4 -f data/ref/chrM.fa -o results/{sample}.vcf results/{sample}.bam
...
```

- Output: ``results/{sample}.vcf`` (uncompressed).
- Guard: ``[[ -f results/{sample}.vcf || -f results/{sample}.vcf.gz ]] || lofreq call-parallel --pp-threads 4 -f data/ref/chrM.fa -o results/{sample}.vcf results/{sample}.bam`` (Check both because step 6 will delete the ``vcf`` and leave ``vcf.gz``.)
- Gotchas: BAM is positional, NOT behind ``-b`/`-i``. The flag is ``--pp-threads``, not ``-t`` or ``--threads``

6. VCF compression + tabix index

```
...
bgzip -f results/{sample}.vcf
...
tabix -p vcf results/{sample}.vcf.gz
...
```

- Outputs: ``results/{sample}.vcf.gz`` and ``results/{sample}.vcf.gz.tbi``.
- Combined guard: ``[[-f results/{sample}.vcf.gz.tbi]] || { bgzip -f results/{sample}.vcf && tabix -p vcf results/{sample}.vcf.gz ; }`
- Gotchas: ``bgzip`` operates ****in place**** - it deletes ``results/{sample}.vcf`` after writing ``vcf.gz``.

7. Collapsed TSV (rebuild every run)

Do NOT guard this step - always overwrite, since per-sample VCFs may have changed.

Header (overwrite):

```
...
printf 'sample\tchrom\tpos\tref\talt\tmaf\n' > results/collapsed.tsv
...
```

Per sample, append:

```
...
```

```
bcftools query -f '%CHROM\t%POS\t%REF\t%ALT\t%INFO/AF\n' results/{sample}.vcf.gz | awk -v s={sample} 'B
...
```

- Gotchas:

- The format string uses `%INFO/AF`, not `%AF` - bcftools requires the `%INFO/` prefix for INFO fields
- The `\t` and `\n` inside `-f '... '` are **bcftools format codes**, parsed by bcftools itself; keep
- awk's `OFS="\t"` is required so `print s,\$0` joins with a tab (`\$0` already contains the tabbed bcf
- Use `>` for the header line, `>>` for every per-sample append.

```
---
```

Idempotency summary

- Steps 1-6 each have a `[[ -f <sentinel> ]] ||` guard on their final output. A second invocation on a
- Step 7 is intentionally rebuilt from scratch on every run (header `>`, then append per sample). This

Plan v2_defensive — plan/PLAN_v2_defensive.md

```
# `run.sh` Implementation Plan - chrM amplicon variant calling
```

0. Script preamble (top of file)

1. First line after shebang: `set -euo pipefail`.
2. Constants: `THREADS=4` and `SAMPLES=("M117-bl" "M117-ch" "M117C1-bl" "M117C1-ch")`.
3. Paths: `REF=data/ref/chrM.fa`, `OUT=results`.
4. `mkdir -p "\$OUT"`.
5. Initialize the failure log every run (truncate, no header - pure TSV body): `: > "\$OUT/failures.log"
6. Track survivors with an array: `SURVIVORS=()` and a counter `OK=0`.

0a. Defensive helper `try`

Define this function exactly:

```
...
```

```
try() { # try <sample> <step_label> <validation_cmd_string> -- <cmd...>
  local sample="$1" step="$2" validate="$3"; shift 3
  [[ "$1" == "--" ]] && shift
  if "$@" && eval "$validate"; then return 0; fi
  "$@" && eval "$validate" && return 0
  printf '%s\t%s\t%s\n' "$sample" "$step" "command_or_validation_failed" >> "$OUT/failures.log"
  return 1
}
...
```

Behavior: runs `cmd`, then evaluates the validation string; on any failure retries the **same cmd + sam

For reference-prep steps (no sample), use sample label `\_\_ref\_\_` and `exit 1` instead of `continue` on

```
---
```

1. Reference preparation (once, before the sample loop)

```
### 1a. `bwa index`
```

```

...
bwa index data/ref/chrM.fa
...

- Outputs: `data/ref/chrM.fa.{amb,ann,bwt,pac,sa}`.
- Idempotency guard: `[[ -f data/ref/chrM.fa.bwt ]] || try __ref__ bwa_index '[[ -s data/ref/chrM.fa.bw
- On failure after retry: `echo "[run.sh] reference index failed" >&2; exit 1`.

### 1b. `samtools faidx`

...
samtools faidx data/ref/chrM.fa
...

- Output: `data/ref/chrM.fa.fai`.
- Guard: `[[ -f data/ref/chrM.fa.fai ]] || try __ref__ faidx '[[ -s data/ref/chrM.fa.fai ]]` -- samtool
- On failure: `exit 1`.

---

## 2. Per-sample loop

`for s in "${SAMPLES[@]}"; do ... done`. Inside the loop, every `try` failure must `continue` to the ne

### Step 2a - Align + sort → `results/{s}.bam`

Exact pipeline (the RG string must be a literal double-quoted string containing the four characters `\\`

...
bwa mem -t 4 -R "@RG\tID:${s}\tSM:${s}\tLB:${s}\tPL:ILLUMINA" \
  data/ref/chrM.fa data/raw/${s}_1.fq.gz data/raw/${s}_2.fq.gz \
  | samtools sort -@ 4 -o results/${s}.bam -
...

- Output: `results/${s}.bam`.
- Guard: `[[ -f results/${s}.bam ]] && samtools quickcheck results/${s}.bam` → skip; else run via `try`
- Validation string passed to `try`: `'samtools quickcheck results/"$s".bam'`.
- Wrap the whole pipeline in a tiny inline shell function or `bash -c` because `try` takes argv, not a
- Gotchas: literal `\\t` only; PL is `ILLUMINA` (uppercase); `samtools sort`'s trailing `-` reads stdin.

### Step 2b - BAM index → `results/{s}.bam.bai`

...
samtools index -@ 4 results/${s}.bam
...

- Output: `results/${s}.bam.bai`.
- Guard: `[[ -s results/${s}.bam.bai ]]` → skip.
- Validation: `'[[ -s results/"$s".bam.bai ]]'`.
- No duplicate marking (amplicon data).

### Step 2c - Variant calling → `results/{s}.vcf`

```

Exact command (no substitutions, no extra flags):

```
...  
lofreq call-parallel --pp-threads 4 -f data/ref/chrM.fa -o results/${s}.vcf results/${s}.bam  
...
```

- Output: `results/${s}.vcf` (uncompressed; lofreq writes plain VCF here).
- Guard: skip if `results/${s}.vcf.gz` already exists and tabix index validates (step 2d covers it); otherwise fail.
- Validation: `'[[ -s results/"$s".vcf ]] && bcftools view -h results/"$s".vcf > /dev/null'`.
- Gotcha: `--pp-threads` is mandatory for `call-parallel`; `-t` is wrong here.

Step 2d - Compress + tabix → `results/${s}.vcf.gz` + `.tbi`

Two separate invocations, each wrapped in its own `try`:

```
...  
bgzip -f results/${s}.vcf  
tabix -p vcf results/${s}.vcf.gz  
...
```

- Outputs: `results/${s}.vcf.gz`, `results/${s}.vcf.gz.tbi`.
- Guard for the pair: `'[[ -s results/${s}.vcf.gz && -s results/${s}.vcf.gz.tbi ]] && bcftools view -h results/${s}.vcf.gz > /dev/null'`.
- Validation after `bgzip`: `'[[ -s results/"$s".vcf.gz ]] && bcftools view -h results/"$s".vcf.gz > /dev/null'`.
- Validation after `tabix`: `'[[ -s results/"$s".vcf.gz.tbi ]]'`.
- Gotcha: `bgzip -f` **deletes** `results/${s}.vcf` on success - that is expected; do not look for it a

Step 2e - Mark survivor

After all four steps succeed for sample `s`: `SURVIVORS+=("$s"); OK=$((OK+1))`.

3. Collapsed TSV (after the loop, only over `SURVIVORS`)

1. Write header (always, overwrite):

```
...  
printf 'sample\tchrom\tpos\tref\talt\taf\n' > results/collapsed.tsv  
...
```

2. For each `s` in `"${SURVIVORS[@]}"`:

```
...  
bcftools query -f '%CHROM\t%POS\t%REF\t%ALT\t%INFO/AF\n' results/${s}.vcf.gz \  
| awk -v s=${s} 'BEGIN{OFS="\t"}{print s,$0}' >> results/collapsed.tsv  
...
```

Wrap in `try "$s" collapse '[[ -s results/collapsed.tsv ]]' -- ...` so a single broken VCF only loses t

4. Final summary + exit code

1. Build a comma-separated list of failed samples by `cut -f1 results/failures.log | sort -u | grep -v`

2. Identify the first failing step per failed sample (e.g. ``awk -F'\t' '!seen[$1]++{print $1" failed at`
3. Emit on `**stderr**`, as the very last line:

...

```
[run.sh] <OK>/<TOTAL> samples completed; <sample> failed at step <label> - see results/failures.log
```

...

If ``OK == ${#SAMPLES[@]}``: ``[run.sh] 4/4 samples completed; no failures``.

4. Exit policy: ``if (( OK >= 1 )); then exit 0; else exit 1; fi``.

5. Idempotency note

Every step is guarded by an ``[[ -f ... ]] && validation`` check before invoking ``try``. A second run on a